



# CS 5331 Special Problems in Computer Science: Applied Data Science

## Gathering and Processing Data --Model Evaluation and Refinement

Dr. Song Liao  
Assistant Professor

September 13, 2024



# Model Evaluation





# Model Evaluation

- In-sample evaluation tells us how well our model will fit the data used to train it
- Problem?
  - It doesn't tell us how well the trained model can be used to predict new data
- Solution
  - In-sample data or training data
  - Out-of-sample evaluation or test set



# Training/Testing set

- Split dataset into:
  - Training set (70%): 
  - Testing set(30%): 
- Build and train the model with a training set
- Use testing set to assess the performance of a predictive model



# Function train\_test\_split()

```
y_data = df['price']    x_data=df.drop('price',axis=1)
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.3, random_state=1)
```

- x\_data: features or independent variables
- y\_data: dataset target: df['price']
- x\_train, y\_train: parts of available data as training set
- x\_test, y\_test: parts of available data as testing set
- test\_size: percentage of the data for testing (here 30%)



# Function train\_test\_split()

```
lre.fit(x_train[['horsepower']], y_train)
```

```
lre.score(x_test[['horsepower']], y_test)
```


0.6287485044222195

```
lre.score(x_train[['horsepower']], y_train)
```

0.6628063120233265

# Cross Validation

- Most common out-of-sample evaluation metrics
- Most effective use of data (each observation is used for both training and testing)





# Cross Validation

```
from sklearn.model_selection import cross_val_score  
  
Rcross = cross_val_score(lre, x_data[['horsepower']], y_data, cv=4)  
  
Rcross  
array([0.7746232 , 0.51716687, 0.74785353, 0.04839605])  
  
print("The mean of the folds are", Rcross.mean(), "and the standard deviation is" , Rcross.std())
```

The mean of the folds are 0.522009915042119 and the standard deviation is 0.291183944475603



# Cross Validation

- Function `cross_val_predict()` returns the prediction that was obtained for each element when it was in the test set

```
from sklearn.model_selection import cross_val_predict
```

```
yhat = cross_val_predict(lre,x_data[['horsepower']], y_data,cv=4)  
yhat[0:5]
```

```
array([14141.63807508, 14141.63807508, 20814.29423473, 12745.03562306,  
       14762.35027598])
```



# Overfitting



```
lr = LinearRegression()
lr.fit(x_train[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_train)
```

```
yhat_train = lr.predict(x_train[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']])
yhat_train[0:5]
```

```
Title = 'Distribution Plot of Predicted Value Using Training Data vs Training Data Distribution'
DistributionPlot(y_train, yhat_train, "Actual Values (Train)", "Predicted Values (Train)", Title)
```




# Overfitting



# Overfitting

## Model Selection



# Overfitting


```
pr = PolynomialFeatures(degree=5)
x_train_pr = pr.fit_transform(x_train[['horsepower']])
x_test_pr = pr.fit_transform(x_test[['horsepower']])
pr
```

```
poly.score(x_train_pr, y_train)
```

```
0.5567716902120254
```


```
poly.score(x_test_pr, y_test)
```

```
-29.871340302043684
```











# Overfitting



# Overfitting






# Ridge Regression

- Ridge regression is a type of linear regression that adds a regularization term to **prevent overfitting** and improve the generalization of the model.


$$\hat{y} = 1 + 2x - 3x^2 - 2x^3 - 12x^4 - 40x^5 + 80x^6 + 71x^7 - 141x^8 - 38x^9 + 75x^{10}$$






# Ridge Regression

```
from sklearn.linear_model import Ridge  
  
RigeModel=Ridge(alpha=0.1)  
  
RigeModel.fit(x_train_pr, y_train)  
  
Rsqu_test = []  
Rsqu_train = []  
dummy1 = []  
ALFA = 10 * np.array(range(0,1000))  
for alfa in ALFA:  
    RigeModel = Ridge(alpha=alfa)  
    RigeModel.fit(x_train_pr, y_train)  
    Rsqu_test.append(RigeModel.score(x_test_pr, y_test))  
    Rsqu_train.append(RigeModel.score(x_train_pr, y_train))
```



# Hyperparameters

- The term alpha in Ridge regression is called hyperparameter
- Scikit-learn has a means of automatically iterating over these hyperparameters using cross-validation called Grid Search





# Grid Search

```
from sklearn.model_selection import GridSearchCV

parameters1= [{'alpha': [0.001,0.1,1, 10, 100, 1000, 10000, 100000, 100000]}]
parameters1

RR=Ridge()

Grid1 = GridSearchCV(RR, parameters1, cv=4)

Grid1.fit(x_data[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_data)

BestRR=Grid1.best_estimator_
BestRR

▼      Ridge
Ridge(alpha=10000)                                BestRR.score(x_test[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_test)
0.8411649831036151
```



# Copy Right Notice

- Most slides in this presentation are adopted from various sources. The Copyright belong to the original authors. Thanks!